Mail to UGVCL Resource Person

Draft & question bank · VidyutOS · PS000125 (UGVCL)

To: hmtcotech2@ugvcl.com · **Send:** after locking PS000125 on the portal, by ~20 Oct 2026 (one mail, from the team-lead SPOC only) · **Follow-up:** one polite reminder after 7 days, then stop. Every answer received must be reflected in `REQUIREMENT_MAPPING.md` (replace assumptions with confirmed facts, citing "per UGVCL resource person, <date>").

Why mail at all

The build currently assumes: SLA timelines (24/48/72/120 h by severity), complaint categories, a simplified org chain (officer -> engineer -> admin), two zones, and synthetic scale figures. Six short answers from UGVCL retire most of that risk before the synopsis (6 Nov) and the finale defence (2 Dec).

Ready-to-send draft (fill the [brackets])

Subject: Seva First Innovation Challenge — PS000125: Student team seeking brief guidance on scope & validation

Respected Sir/Madam,

We are [Team Name] from [Institute], participating in the Seva First Innovation Challenge. We have selected Problem Statement **PS000125** (Future-Ready, Secure and Intelligent IT Infrastructure & User-Friendly ERP/Portal Ecosystem) and are building a working prototype: a unified role-based portal for new connections and complaints with automated workflow routing, SLA tracking, an AI assistant over policy documents, and offline field-work sync for low-connectivity areas. We use **synthetic data only** and are not seeking any production access or real consumer data.

To align our synopsis (due 6 Nov) with UGVCL's actual priorities, we would be grateful for brief guidance on:

1. Which of the problem areas (duplicate data entry, portal integration, offline field work, SLA follow-up, dashboards) is the highest priority for UGVCL today?
2. The typical steps and official timelines/SLAs for a new connection request, and the complaint categories used in your current system.
3. The approval hierarchy (Corporate/Circle/Division/Subdivision) a request passes through, and which staff roles handle it at each level.
4. Field conditions: devices used by field staff and areas with poor connectivity, so our offline design is realistic.
5. Any public SOPs/circulars (e.g., new-connection procedure) we may reference, and any sample *formats* (blank forms/report layouts) you can share — structure only, no real data.
6. For the infrastructure part of the statement: is a phased modernization *proposal* expected, and are there hosting/security preferences (cloud/on-premise, SSO) we should design for?
7. Approximate scale we should design for (subdivisions, consumers, requests per day).

Written replies to any subset, or a short 20-minute call at your convenience, would help us greatly. We would also welcome the opportunity to show the prototype to your team before the Grand Finale (2 Dec) if that is feasible.

Thank you for your time and for posing this problem statement.

[Team Lead Name] · [Team Name] · [Institute] · [Phone] · [Email]

Question bank (if a call happens — pick in this order)

Validate the build (highest value)

- [ ] Priority ranking of the 12 problem points
- [ ] Real SLA hours per category/severity vs our assumed 24/48/72/120
- [ ] Official complaint/request taxonomy vs our ML categories
- [ ] Approval chain & role names (do "officer/engineer" map to JE/DE/EE?)
- [ ] Scale: subdivisions, daily request volume, peak load

Resources that would materially help

- [ ] Blank/redacted forms or report formats (structure only)
- [ ] Public SOPs/circulars the assistant may quote

- [ ] List of current ERP modules/portals + any existing API layer
- [ ] Field-device reality (Android? BYOD? worst-connectivity areas)
- [ ] Language needs: Gujarati UI needed for consumers/field staff?

Scope & expectations

- [ ] Hardware point: proposal-only or costing expected? Existing standards?
- [ ] Hosting/security preferences: cloud vs on-prem, SSO, DPDP/audit needs
- [ ] Secondary users: do contractors/agencies or govt departments need logins in scope?
- [ ] Finale expectations: live prototype vs slides; pilot subdivision possible?
- [ ] Any visit/demo opportunity before 2 Dec

Never ask / never do

- No real consumer data, credentials, or production access (state synthetic-only up front — done in the draft).
- No questions answerable from public sources (tariffs, company profile).
- No repeat nudges beyond one reminder; no mailing other UGVCL addresses found online.
- Don't promise features in the mail that are [Roadmap] roadmap in `REQUIREMENT_MAPPING.md`.

After replies arrive

1. Update `REQUIREMENT_MAPPING.md` statuses/notes with confirmed facts + date.
2. If SLA hours/categories differ, Role 1 (backend) + Role 3 (ML) adjust `workflow.py`/taxonomy and retrain (`make train`) — this is why the contract freeze is 24 Oct.
3. Quote *one* confirmed fact in the synopsis problem section ("per UGVCL…") — judges reward validated problems.

Plan A — if UGVCL never replies (role-wise: fully doable without their help)

The prototype already runs end-to-end on synthetic data. No role is blocked.

Role	Can finish without UGVCL	How the assumption risk is contained
R1 Backend & Workflow	Entire workflow, routing, SLA clocks, audit, all APIs, seed data	Keep SLA hours, categories and zones as **config values** (`config.py`), so a late correction is a one-line change, not a refactor
R2 Frontend & PWA	All four portals, offline queue/sync, notifications, reports	Role labels stay generic ("Officer/Engineer"); an i18n-ready structure means Gujarati strings can drop in later without layout changes
R3 ML & AI	Both models trained on labelled synthetic data; assistant answers from clearly-labelled demo policy extracts	Taxonomy lives in one list — relabelling = edit + `make train` (minutes, not days)
R4 DevOps/Security/QA	Docker, Makefile, backups, tests, security defaults, audit pack	Defaults follow public best practice (JWT/RBAC/bcrypt/CORS); any UGVCL-specific policy would be config/SSO work, not redesign
R5 Docs/Demo/Product	Synopsis, deck, demo, mapping docs — complete today	Every assumption is **labelled as an assumption** in `REQUIREMENT_MAPPING.md`; honest framing scores better than silent guessing

Rule: after the 24 Oct contract freeze, no role waits for the mail. The demo on 2 Dec works with zero UGVCL input.

Plan B — after resources/replies arrive (role-wise: plug-in upgrades)

Treat every reply as a drop-in replacement for an assumption — budget ~1–2 days total, none of it structural.

Role	Upgrade unlocked	Effort
R1 Backend & Workflow	Real SLA matrix + approval hierarchy (JE/DE/EE titles, circle/division tree) into config + seed roles	Config/seed edit + tests still green

Role	Upgrade unlocked	Effort
R2 Frontend & PWA	Real form fields & report formats; official role names; Gujarati labels if UGVCL wants them	Label/field swaps on existing screens
R3 ML & AI	Retrain on the **official complaint taxonomy**; replace demo policy extracts with real public SOPs/circulars in the assistant (RAG source swap)	Same pipeline — `make train`, documents folder swap
R4 DevOps/Security/QA	Align to stated hosting (cloud/on-prem), SSO, DPDP/audit-retention expectations; size infra tiers with real scale numbers	Settings + `INFRASTRUCTURE_ROADMAP.md` update with confirmed figures
R5 Docs/Demo/Product	Synopsis quotes confirmed facts ("per UGVCL resource person…"); mapping rows flip from *assumption* to *validated*; if offered, schedule the pre-finale demo/visit	Doc edits + demo-script tweak; biggest marks gain per hour spent

Deadline logic: replies before **24 Oct** -> fold into the synopsis draft directly. Replies **after 24 Oct** -> apply only config/label swaps (Plans above), nothing structural, per the freeze. No reply at all -> Plan A, and the mail itself is still evidence of stakeholder outreach in the deck.